

Agreement Protocols

Classification of Faults

- Based on components that failed
 - Program / process
 - Processor / machine
 - Link
 - Storage
 - Clock
- Based on behavior of faulty component
 - Crash – just halts
 - Failstop – crash with additional conditions
 - Omission – fails to perform some steps
 - Byzantine – behaves arbitrarily
 - Timing – violates timing constraints

Classification of Tolerance

- Types of tolerance:
 - Masking – system always behaves as per specifications even in presence of faults
 - Non-masking – system may violate specifications in presence of faults. Should at least behave in a well-defined manner
- Fault tolerant system should specify:
 - Class of faults tolerated
 - What tolerance is given from each class

Core problems

- Agreement (multiple processes agree on some value)
- Clock synchronization
- Stable storage (data accessible after crash)
- Reliable communication (point-to-point, broadcast, multicast)
- Atomic actions

Overview of Consensus Results

- Let f be the maximum number of faulty processors.
- Tight bounds for message passing:

| | Crash failures | Byzantine failures |
|----------------------------|----------------|--------------------|
| Number of rounds | $f + 1$ | $f + 1$ |
| Total number of processors | $f + 1$ | $3f + 1$ |
| Message size | polynomial | polynomial |

Overview of Consensus Results

- *Impossible* in asynchronous case.
 - Even if we only want to tolerate a single crash failure.
 - True both for message passing and shared read-write memory.

Consensus Algorithm for Crash Failures

Code for each processor:

$v :=$ my input

at each round 1 through $f+1$:

if I have not yet sent v then send v to all

wait to receive messages for this round

$v :=$ minimum among all received values and

current value of v

if this is round $f+1$ then decide on v


Correctness of Crash Consensus Algo

- Termination: By the code, finish in round $f + 1$.
- Validity: Holds since processors do not introduce spurious messages
 - if all inputs are the same, then that is the only value ever in circulation.

Correctness of Crash Consensus Algo

Agreement:

- Suppose in contradiction p_j decides on a smaller value, x , than does p_i .
- Then x was hidden from p_i by a chain of faulty processors:



- There are $f + 1$ faulty processors in this chain, a contradiction.

Performance of Crash Consensus Algo

- Number of processors $n > f$
- $f + 1$ rounds
- $n^2 \cdot |V|$ messages, each of size $\log |V|$ bits, where V is the input set.

Lower Bound on Rounds

Assumptions:

- $n > f + 1$
- every processor is supposed to send a message to every other processor in every round
- Input set is $\{0,1\}$

Byzantine Agreement Problems

Model :

- Total of n processes, at most m of which can be faulty
- Reliable communication medium
- Fully connected
- Receiver always knows the identity of the sender of a message
- Byzantine faults
- Synchronous system
 - In each round, a process receives messages, performs computation, and sends messages.

Byzantine Agreement

- Also known as Byzantine Generals problem
 - One process x broadcasts a value v
 - Agreement Condition: All non-faulty processes must agree on a common value.
 - Validity Condition: The agreed upon value must be v if x is non-faulty.

Variants


- Consensus
 - Each process broadcasts its initial value
 - Satisfy agreement condition
 - If initial value of all non-faulty processes is v , then the agreed upon value must be v
- Interactive Consistency
 - Each process k broadcasts its own value v_k
 - All non-faulty processes agree on a common vector (v_1, v_2, \dots, v_n)
 - If the k^{th} process is non-faulty, then the k^{th} value in the vector agreed upon by non-faulty processes must be v_k
- *Solution to Byzantine agreement problem implies solution to other two*

Byzantine Agreement Problem


- No solution possible if:
 - asynchronous system, or
 - $n < (3m + 1)$
- Lower Bound:
 - Needs at least $(m+1)$ rounds of message exchanges
- “*Oral*” messages – messages can be forged / changed in any manner, but the receiver always knows the sender

Proof


Theorem: *There is no t -Byzantine-robust broadcast protocol for $t \geq N/3$*



Scenario-0: **T must decide 0**



Scenario-1: **U must decide 1**



Scenario-2:

- similar to Scenario-0 for T
- similar to Scenario-1 for U
- T decides 0 and U decides 1

Lamport-Shostak-Pease Algorithm

- Algorithm *Broadcast*(N, t) where t is the resilience

For $t = 0$, *Broadcast*($N, 0$):

Pulse

- 1 The general sends $\langle \text{value}, x_g \rangle$ to all processes,
the lieutenants do not send.

Receive messages of pulse 1.

The general decides on x_g .

Lieutenants decide as follows:

if a message $\langle \text{value}, x \rangle$ was received from g in pulse-1
then decide on x
else decide on *undef*

Lamport-Shostak-Pease Algorithm contd..

For $t > 0$, $Broadcast(N, t)$:

Pulse

- 1 The general sends $\langle \text{value}, x_g \rangle$ to all processes, the lieutenants do not send.
Receive messages of pulse 1.
Lieutenant p acts as follows:
 - if a message $\langle \text{value}, x \rangle$ was received from g in pulse-1 then $x_p = x$ else $x_p = \text{undef}$;
 - Announce x_p to the other lieutenants by acting as a general in $Broadcast_p(N - 1, t - 1)$ in the next pulse

Pulse

- $t + 1$ Receive messages of pulse $t + 1$.
The general decides on x_g .
For lieutenant p :
 - A decision occurs in $Broadcast_q(N - 1, t - 1)$ for each lieutenant q
 - $W_p[q] = \text{decision in } Broadcast_q(N - 1, t - 1)$
 - $y_p = \text{major}(W_p)$

Features

- Termination: If $Broadcast(N, t)$ is started in pulse 1, every process decides in pulse $t + 1$
- Dependence: If the general is correct, if there are f faulty processes, and if $N > 2f + t$, then all correct processes decide on the input of the general
- Agreement: All correct processes decide on the same value

The Broadcast(N, t) protocol is a t -Byzantine-robust broadcast protocol for $t < N/3$

Time complexity: $O(t + 1)$ Message complexity: $O(N^t)$